

Terraform 101

A talk by Charles

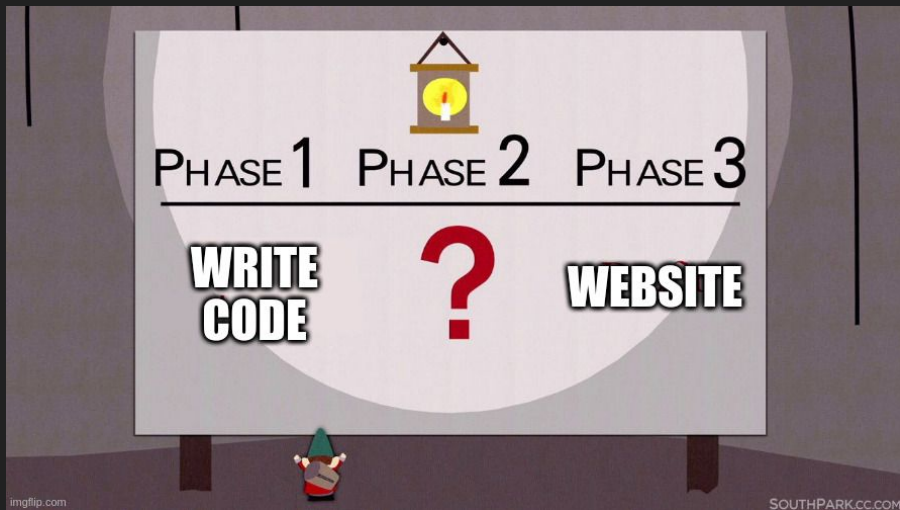
Slides and code: <https://github.com/bushong1/presentations/tree/main/terraform-101>

Agenda

- What's Infrastructure?
- Terraform Intro and Anatomy
- Actually running it
- Good Stuff
- Helpful Tips + Brain Dump

Infrastructure?

Infrastructure is...?



- **Everything ?** could be “Infrastructure”
 - Automated tests
 - Static code to bucket
 - Docker image build
 - Container Deploy
 - Server Build
 - Database
 - Caching
 - Network
 - Security
 -

Let's make a server on AWS

How do... computer?

- To create a server on AWS
 - a. Know all the answers to the values on the right
 - b. Pray you didn't get it wrong
 - c. Get hacked because you did it wrong

```
"AdditionalInfo" : String,  
"Affinity" : String,  
"AvailabilityZone" : String,  
"BlockDeviceMappings" : [ BlockDeviceMapping, ... ],  
"CpuOptions" : CpuOptions,  
"CreditSpecification" : CreditSpecification,  
"DisableApiTermination" : Boolean,  
"EbsOptimized" : Boolean,  
"ElasticGpuSpecifications" : [ ElasticGpuSpecification, ... ],  
"ElasticInferenceAccelerators" : [ ElasticInferenceAccelerator, ... ],  
"EnclaveOptions" : EnclaveOptions,  
"HibernationOptions" : HibernationOptions,  
"HostId" : String,  
"HostResourceGroupArn" : String,  
"IamInstanceProfile" : String,  
"ImageId" : String,  
"InstanceInitiatedShutdownBehavior" : String,  
"InstanceType" : String,  
"Ipv6AddressCount" : Integer,  
"Ipv6Addresses" : [ InstanceIpv6Address, ... ],  
"KernelId" : String,  
"KeyName" : String,  
"LaunchTemplate" : LaunchTemplateSpecification,  
"LicenseSpecifications" : [ LicenseSpecification, ... ],  
"Monitoring" : Boolean,  
"NetworkInterfaces" : [ NetworkInterface, ... ],  
"PlacementGroupName" : String,  
"PrivateIpAddress" : String,  
"RamdiskId" : String,  
"SecurityGroupIds" : [ String, ... ],  
"SecurityGroups" : [ String, ... ],  
"SourceDestCheck" : Boolean,  
"SsmAssociations" : [ SsmAssociation, ... ],  
"SubnetId" : String,  
"Tags" : [ Tag, ... ],  
"Tenancy" : String,  
"UserData" : String,  
"Volumes" : [ Volume, ... ]
```


But maybe you suffer through it

...once

What if I wanna do it again?

- That's hard, and I forget stuff...
 - Write documentation
- What if it changes?
 - Update the docs
- You won't.
 - This is hard?

Enter: AWS API



But there's a lot of API...

accessanalyzer	cloudwatch	ds	history	location	pinpoint-email	ses
acm	codeartifact	dynamodb	honeycode	logs	pinpoint-sms-voice	sesv2
acm-pca	codebuild	dynamodbstreams	iam	lookoutvision	polly	shield
alexaforbusiness	codecommit	ebs	identitystore	machinelearning	pricing	signer
amp	codeguruprofiler	ec2	imagebuilder	macie	qldb	sms
amplify	codeguru-reviewer	ec2-instance-connect	importexport	macie2	qldb-session	snowball
amplifybackend	codepipeline	ecr	inspector	managedblockchain	quicksight	sns
apigateway	codestar	ecr-public	iot	marketplace-catalog	ram	sqs
apigatewaymanagementapi	codestar-connections	ecs	iot1click-devices	marketplace-commerceanalytics	rds	ssm
apigatewayv2	codestar-notifications	efs	iot1click-projects	marketplace-entitlement	rds-data	sso
appconfig	cognito-identity	eks	iotanalytics	mediaconnect	redshift	redshift
appflow	cognito-idp	elasticache	iot-data	mediaconvert	redshift-data	sso-oidc
appintegrations	cognito-sync	elasticbeanstalk	iotdeviceadvisor	medialive	rekognition	stepfunctions
application-autoscaling	comprehend	elastic-inference	iotevents	mediapackage	resource-groups	storagegateway
application-insights	comprehendmedical	elastictranscoder	iotevents-data	mediapackage-vod	resourcegroupstaggingapi	sts
appmesh	compute-optimizer	elb	iotfleethub	mediastore	robomaker	support
appstream	configservice	elbv2	iot-jobs-data	mediastore-data	route53	swf
appsync	configure	emr	iotsecuretunneling	mediatailor	route53domains	synthetics
athena	connect	emr-containers	iotsitewise	meteringmarketplace	route53resolver	texttract
auditmanager	connect-contact-lens	es	iotthingsgraph	mgh	s3	timestream-query
autoscaling	connectparticipant	events	iotwireless	migrationhub-config	s3api	timestream-write
autoscaling-plans	cur	firehose	ivs	mobile	s3control	transcribe
backup	customer-profiles	fms	kafka	mq	s3outposts	transfer
batch	databrew	forecast	kendra	mturk	sagemaker	translate
braket	dataexchange	forecastquery	kinesis	mwaa	sagemaker-a2i-runtime	waf
budgets	datapipeline	frauddetector	kinesisanalytics	neptune	sagemaker-edge	waf-regional
ce	datasync	fsx	kinesisanalyticsv2	network-firewall	sagemaker-featurestore-runtime	wafv2
chime	dax	gamelift	kinesisvideo	networkmanager	sagemaker-runtime	wellarchitected
cli-dev	ddb	glacier	kinesis-video-archived-media	opsworks	savingsplans	workdocs
cloud9	deploy	globalaccelerator	kinesis-video-media	opsworks-cm	schemas	worklink
clouddirectory	detective	glue	kinesis-video-signaling	opsworkscm	sdb	workmail
cloudformation	devicefarm	greengrass	kms	organizations	secretsmanager	workmailmessageflow
cloudfront	devops-guru	greengrassv2	lakeformation	outposts	securityhub	workspaces
cloudhsm	directconnect	groundstation	lambda	personalize	serverlessrepo	xray
cloudhsmv2	discovery	guardduty	lex-models	personalize-events	servicecatalog	
cloudsearch	dlm	health	lex-runtime	personalize-runtime	servicecatalog-appregistry	
cloudsearchdomain	dms	healthlake	license-manager	pi	servicediscovery	
cloudtrail	docdb	help	lightsail	pinpoint	service-quotas	

Creating your infrastructure could be done...

- by hand in the GUI
 - Easy to start, hard to maintain
- by hand with the CLI
 - Easy to forget how
- with a homemade script using the API
 - Hard to link stuff together, track state
- **using a tool purpose made to solve these problems**



Terraform

- API Wrapper
 - It does what AWS lets it
- State Manager
 - “What do I have? What do I need? What changed?”
- Modular, Declarative Language
 - YAML-esque
- Fast
 - ...-er than alternatives

Let's get to the code

Anatomy of Terraform

```
locals {  
  app_name = "hello-world"  
}  
data "aws_ami" "ubuntu" {  
  most_recent = true  
  filter {  
    name     = "name"  
    values   = ["ubuntu-20.04-<...>-*"]  
  }  
  filter {  
    name     = "virtualization-type"  
    values   = ["hvm"]  
  }  
  owners = ["099720109477"] # Canonical  
}  
  
resource "aws_instance" "my_server" {  
  ami             = data.aws_ami.ubuntu.id  
  instance_type   = "t3.micro"  
  
  tags = {  
    Name = local.app_name  
  }  
}
```

Anatomy of Terraform

- **Resource**
 - Creates a “Thing”

```
locals {
  app_name = "hello-world"
}
data "aws_ami" "ubuntu" {
  most_recent = true
  filter {
    name      = "name"
    values    = ["ubuntu-20.04-<...>-*"]
  }
  filter {
    name      = "virtualization-type"
    values    = ["hvm"]
  }
  owners = ["099720109477"] # Canonical
}

resource "aws_instance" "my_server" {
  ami           = data.aws_ami.ubuntu.id
  instance_type = "t3.micro"

  tags = {
    Name = local.app_name
  }
}
```


Anatomy of Terraform

- **Resource**
 - Creates a “Thing”
- **Data Source**
 - Loads a “Thing”

```
locals {
  app_name = "hello-world"
}

data "aws_ami" "ubuntu" {
  most_recent = true
  filter {
    name      = "name"
    values    = ["ubuntu-20.04-<...>-*"]
  }
  filter {
    name      = "virtualization-type"
    values    = ["hvm"]
  }
  owners = ["099720109477"] # Canonical
}

resource "aws_instance" "my_server" {
  ami          = data.aws_ami.ubuntu.id
  instance_type = "t3.micro"

  tags = {
    Name = local.app_name
  }
}
```

Anatomy of Terraform

- **Resource**
 - Creates a “Thing”
- **Data Source**
 - Loads a “Thing”
- **Local**
 - A private variable

```
locals {  
  app_name = "hello-world"  
}  
  
data "aws_ami" "ubuntu" {  
  most_recent = true  
  filter {  
    name     = "name"  
    values   = ["ubuntu-20.04-<...>-*"]  
  }  
  filter {  
    name     = "virtualization-type"  
    values   = ["hvm"]  
  }  
  owners = ["099720109477"] # Canonical  
}  
  
resource "aws_instance" "my_server" {  
  ami           = data.aws_ami.ubuntu.id  
  instance_type = "t3.micro"  
  
  tags = {  
    Name = local.app_name  
  }  
}
```

Now what?

- <hand-wave AWS creds>
- `terraform plan`
- List of everything to add, change, or destroy

An execution plan has been generated and is shown below.
Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
# aws_instance.my_server will be created
+ resource "aws_instance" "my_server" {
  + ami                        = "ami-042e8287309f5df03"
  + arn                       = (known after apply)
  + associate_public_ip_address = (known after apply)
  + availability_zone         = (known after apply)
  + cpu_core_count            = (known after apply)
  + cpu_threads_per_core      = (known after apply)
  + get_password_data         = false
  + host_id                   = (known after apply)
  + id                        = (known after apply)
  + instance_state            = (known after apply)
  + instance_type             = "t3.micro"
  + ipv6_address_count        = (known after apply)
```

...

```
  + volume_type              = (known after apply)
}
}
```

Plan: 1 to add, 0 to change, 0 to destroy.

THIS IS HUGE

And then...

- ``terraform apply``
- Does a plan again
- “yes” to apply the changes

```
Plan: 1 to add, 0 to change, 0 to destroy.
```

```
Do you want to perform these actions?
```

```
Terraform will perform the actions described above.
```

```
Only 'yes' will be accepted to approve.
```

```
Enter a value: yes
```

```
aws_instance.my_server: Creating...
```

```
aws_instance.my_server: Still creating... [10s elapsed]
```

```
aws_instance.my_server: Creation complete after 17s [id=i-0840e1e43425b2504]
```

And then...

- ``terraform apply``
- Does a plan again
- “yes” to apply the changes

And then...

- ``terraform apply``
- Does a plan again
- But nothing changes

```
$ terraform apply
aws_instance.my_server: Refreshing state... [id=i-0840e1e43425b2504]

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

And then...

- ``terraform apply``
- Does a plan again
- “yes” to apply the changes

And then...

- ``terraform apply``
- Does a plan again
- But nothing changes

And then...

- ``terraform destroy``
- Does a plan again, deleting stuff
- “yes” to destroy everything

```
Plan: 0 to add, 0 to change, 1 to destroy.
```

```
Do you really want to destroy all resources?
```

```
Terraform will destroy all your managed infrastructure, as shown above.  
There is no undo. Only 'yes' will be accepted to confirm.
```

```
Enter a value: yes
```

```
aws_instance.my_server: Destroying... [id=i-0840e1e43425b2504]
```

```
aws_instance.my_server: Still destroying... [id=i-0840e1e43425b2504, 10s elapsed]
```

Get to the good stuff!

Any number of Resources

- `count` - Create multiple of a thing
- Within that resource, `count.index`

```
resource "aws_instance" "my_server" {  
  count = 5  
  ami           = data.aws_ami.ubuntu.id  
  instance_type = "t3.micro"  
  
  tags = {  
    Name = "${local.app_name}-${count.index}"  
  }  
}
```

Array indexing

- Reference another array using index

```
locals {  
  server_names = [  
    "dog",  
    "cat",  
    "lizard",  
  ]  
}  
  
resource "aws_instance" "my_server_named_better" {  
  count = 3  
  ami           = data.aws_ami.ubuntu.id  
  instance_type = "t3.micro"  
  
  tags = {  
    Name = "my-server-${local.server_names[count.index]}"  
  }  
}
```

Local Modules

- All files in one directory are read together
- Reference files in other directories with modules
- Repeatable

```
locals {
  server_config = {
    dog : { instance_type = "m3.large" }
    cat : { instance_type = "m3.medium" }
    lizard : { instance_type = "m3.small" }
  }
}

module "ami" {
  source = "../ami-lookup"
}

module "my-servers" {
  source      = "../my-server"
  for_each   = local.server_config

  server_name      = "server-${each.key}"
  server_instance_type = each.value.instance_type
  ami_id           = module.ami.ubuntu
}
```

Local Modules - New Resources

- Variable

- Define in your module
- Assign during module creation

```
variable "ami_id" { type = string }  
variable "server_name" { type = string }  
variable "server_instance_type" { type = string }
```

- Output

- Reference with module name
- Output from your module

```
module "my-servers" {  
  source      = "./my-server"  
  for_each    = local.server_config
```

```
  ami_id      = module.ami.ubuntu  
  server_name = "server-${each.key}"  
  server_instance_type = each.value.instance_type
```

```
}
```

```
output "ubuntu" {  
  value = data.aws_ami.ubuntu.id  
}
```

Public Modules

- Terraform Registry
- Some pseudo-officially supported
- YMMV
- Good ones:
 - VPC: <https://registry.terraform.io/modules/terraform-aws-modules/vpc/aws/>
 - Fully featured VPC
 - Handles lots of pitfalls
 - USSBA: <https://registry.terraform.io/search/modules?q=ussba>
 - Managed by us!

Absolutely no way I have time for this,
but here it is just in case
or for reference or whatever

Statefile Management

- Statefiles contain the metadata for your resources
- S3 Backend
- Locking?
- Secrets (RDS!!!!!!)
- Imports?
 - Create a resource outside terraform, then import into existing state

Workspaces

- Parallel sets of resources using the same terraform code
- Use cases:
 - Prod/test/dev
 - Multiple dev environments (per dev or app)
- Each has its own state file
- Base defaults in terraform variables file
- Workspace specific variable values in <workspace>.tfvars file
- `${terraform.workspace}` is a handy way to propagate the env name into tags and resource names

Security

- Any secrets that find their way into the statefile are in the clear
- Tfvars may be a good place to list these, but make sure those don't wind up in Github or a public S3 bucket
- Terraform apply can read AWS creds from environment variables, so you don't have to put your IAM creds in tfvars, regardless of what tutorials say
- Provide non-repudiation by forcing users to use their own IAM creds and assume a "Terraformer" role

Multi-Account

Considerations with two or more *AWS* accounts:

- Scope of named resources
 - Globally Unique? (S3 Bucket)
 - Unique per Account? (Security Group Name)
 - Just a tag? (EC2 Instance “Name”)
- Accidentally running on the wrong account
 - “provider” has an `allowed_account_ids` field

Monolithic vs Multi-state

- Monolith
 - Everything in one Terraform state
 - Good
 - One stop shop
 - Only one place to apply
 - Dependencies easily available
 - Bad
 - Long-running apply
 - Spaghetti code
 - Conflicts resolution
 - Toe Stepping
- Multi-state
 - One state per repeated context
 - Good
 - Easier to read sections
 - Faster individual apply
 - Bad
 - Mutli-state x multi-workspace = confusing
 - “Where did I put that thing?”
 - Dependency Injection

Neat Commands

- terraform <...>
 - **fmt -recursive**
 - Format all your code
 - **validate**
 - Confirm code will (mostly) run properly
 - **show**
 - Eyeball the state
 - **taint <resource>**
 - Mark a resource for recreation
 - **console**
 - Bring up a CLI to play with resources; not perfect